

Design Patterns for Concurrent Objects

Marek Biskup

Warsaw University

Outline

- **Design patterns**
- **Synchronization patterns**
- **Concurrency patterns**
- **Event-driven programming**

Design Patterns



- **Some software design problems reappear in many application developments**
- **Successful solutions to such problems may be reused**
- **Design Patterns – documented reusable solution**
 - Not a finished design that can be transferred into code
 - Idea of how to solve a problem in a particular context
- **Mostly related to Object-Oriented design**
- **Give us common vocabulary to discuss O-O design**

Example – Proxy Pattern

- **Proxy: An object acting as an interface to another object**
 - All methods of the original object but doing something "in between"
- **CacheProxy**
 - Results of expensive operations are cached so that the next call for the same data will return immediately
- **RemoteProxy (e.g. JavaRMI stubs)**
 - Client calls a method of a local stub
 - The stubs sends method parameters to the server
 - The server executes the method
 - Results are returned from the stub as if the call were local

Example – Abstract Factory

- **Decouples creation of objects from the implementation**
- **Example: Bad GUI desing – code dependent on the platform**
 - Windows code:
 - `Button b = new WinButton();`
 - Linux code:
 - `Button b = new QtButton();`
- **AbstractFactory – platform independent code**
 - `Button b = guiFactory.getButton();`
- **Initialization (once in the application)**
 - `guiFactory = new WinGuiFactory();`



Synchronization

- Access to the same memory by several threads has to be synchronized
- Example: increasing a value

```
int value; // global
void increase() {
    int x = value
    x++;
    value = x;
}
```
- Executed on two processors simultaneously may increase the value by one instead of two
- There are standard means of synchronization: mutexes (lock and unlock), monitors (as in java), semaphores

Locking

- **Critical section has to be guarded by a lock**
- **When leaving the critical section the lock must be released, but:**
 - There are different execution paths
 - Exceptions may be thrown
- **Failing to release a lock results in a deadlock**
- **Such an error is difficult to spot**
- **And the code is difficult to maintain**

```
class Counter {  
    bool increment(int) {  
        lock.lock();  
        if ( ... ) {  
            // ...  
            lock.unlock();  
            return false;  
        }  
        // ...  
        lock.unlock();  
        return true;  
    }  
};
```

Scoped Locking Pattern

- **Solution: Use a class that:**
 - Acquires the lock in the constructor
 - Releases the lock in the destructor
- **Then the lock will always be released**
- **Use that pattern even if you have a critical section inside a block:**

```
mutex lock;  
void f() {  
    // do whatever  
    {  
        MutexGuard(lock);  
        // critical section  
        // ...  
    }  
}
```

```
class MutexGuard {  
    Mutex* lock;  
public:  
    MutexGuard(Mutex& aLock)  
        : lock(aLock) {  
        lock->lock();  
    }  
    ~MutexGuard() {  
        lock->unlock();  
    }  
};
```

```
class counter {  
    Mutex lock;  
    bool increment() {  
        MutexGuard(lock);  
        // critical section  
    } // guard will be released in ~counter()  
};
```


Locking in Java

- Use *synchronized* statement
- With an explicit lock: no local objects, so no destructors will be called
 - Cannot use the Scoped Locking Pattern
- Solution: the *finally* statement
- Finally block is always executed when leaving a try block:
 - Normal leaving
 - After the return statement
 - When an exception is thrown

```
synchronized int aFunction() {  
    // synchronized code  
}
```

```
public int aFunction()  
    throws MyException {  
    lock.lock();  
    try {  
        if (A < 0)  
            throw new MyException(A);  
        if (A == 1)  
            return 1;  
        A++;  
        return 0;  
    } finally {  
        lock.unlock();  
    }  
}
```

Intra-Component Call

- Often there is need to call another method from the same component (class)
- If component methods are synchronized the same lock will be acquired twice
 - Deadlock if the lock is not reentrant
 - Unnecessary overhead with reentrant locks

```
class buffer {  
    Mutex lock;  
    public:  
        void put(int a) {  
            MutexGuard(lock);  
            A[num++] = a;  
            if (num == N)  
                flush(); // deadlock!  
        }  
        void flush() {  
            MutexGuard(lock);  
            // write the contents to a file  
        }  
};
```

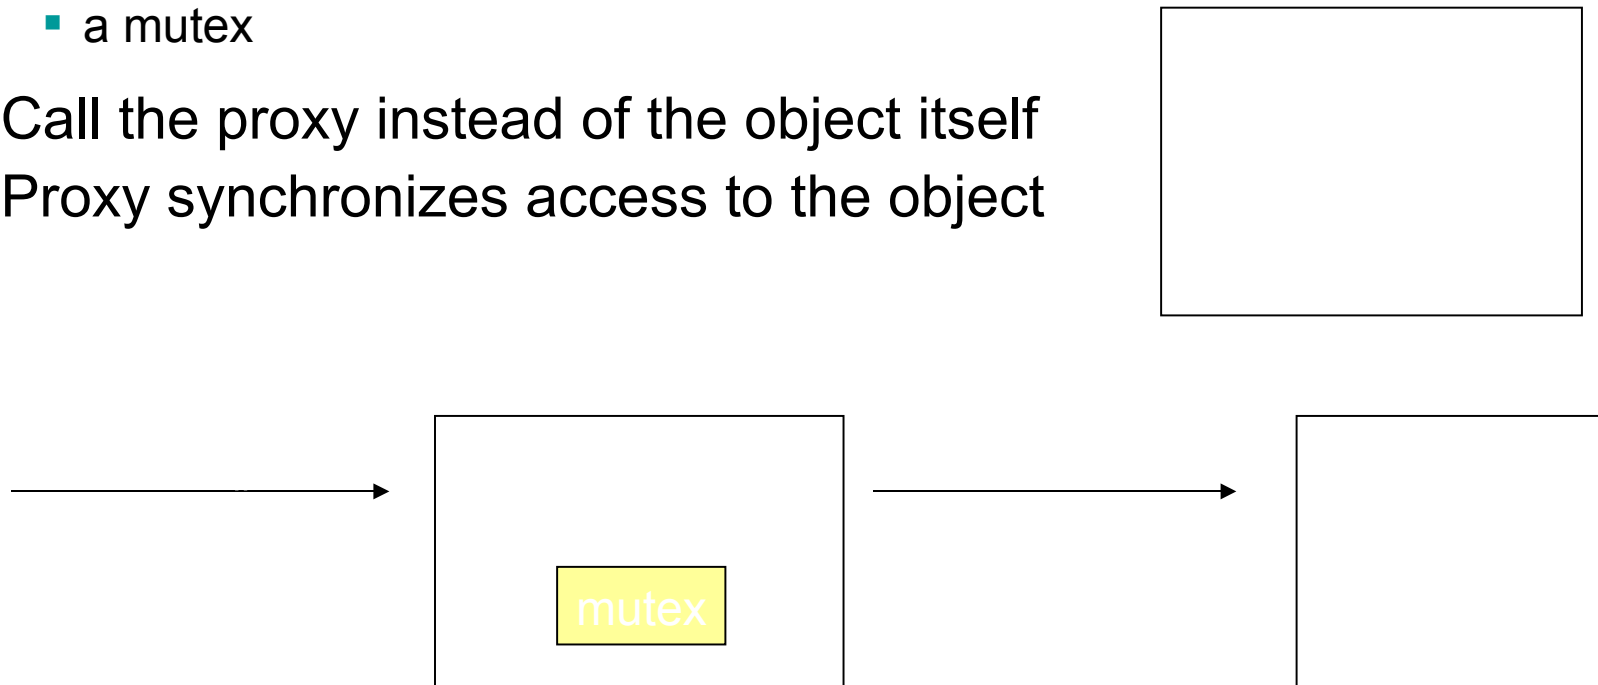
Thread-Safe Interface

- **Structure the component into two layers**
 - Interface methods – acquire a lock and call an implementation method
 - Implementation methods – assume that the lock has already been acquired and do the job
- **Implementation methods may only call other implementation methods (and not interface ones)**
- **But:**
 - Still be careful when calling other component methods!
 - It adds a bit of overhead (more methods, twice more function calls)

```
class buffer {  
    Mutex lock;  
    public:  
        void put(int a) {  
            MutexGuard(lock);    putImp();  
        }  
        void flush() {  
            MutexGuard(lock);    flushImp();  
        }  
    protected:  
        void putImp() {  
            A[num++] = a;  
            if (num == N)  
                flushImp(); // implementation  
        }  
        void flushImp() {  
            // write the contents to a file  
        }  
};
```

Synchronizing Proxy

- **Synchronize access to an object without changing its code**
- **Solution: Synchronizing Proxy Pattern**
 - The proxy contains:
 - all the methods from our object
 - a mutex
 - Call the proxy instead of the object itself
 - Proxy synchronizes access to the object



Read-Write Lock Pattern

- **Consider a concurrent dictionary:**
 - Many threads can read at the same time
 - Exclusive access when writing
- Better locking strategy – a lock with two operations:
 - `readLock();` – blocks when the write-lock is taken
 - `writeLock();` – blocks when the read or write-lock is taken

ReadWriteLock

- Starvation problem:
 - When readers keep on coming, writers will wait forever
- Solution:
 - e.g.: don't allow readers when a writer is waiting

Initializing a value

- Think of a value that should be initialized when accessed for the first time

// no synchronization

```
class Foo {  
    private Helper = null;  
    public Helper getHelper() {  
        if (helper == null)  
            helper = new Helper();  
        return helper;  
    }  
    // other functions and members...  
}
```

// too much synchronization (?)

```
class Foo {  
    private Helper = null;  
    public synchronized Helper getHelper() {  
        if (helper == null)  
            helper = new Helper();  
        return helper;  
    }  
    // other functions and members...  
}
```

Double Checked Locking

■ Can we do it better?

// double checked locking optimization

```
class Foo {  
    private Helper = null;  
    public Helper getHelper() {  
        if (helper == null) {  
            synchronized(this) {  
                if (helper == null) {  
                    helper = new Helper();  
                }  
            }  
            return helper;  
        }  
    }  
}
```

Double check:
several processes
may enter here

Scenario:

Thread A:

```
helper == null;  
helper = new Helper();
```

compilers may assign
the pointer value before
A has finished executing
the constructor!

Thread B:

```
helper != null; //!  
useHelper // not initialized!  
CRASH!
```

D.C.L. is an AntiPattern

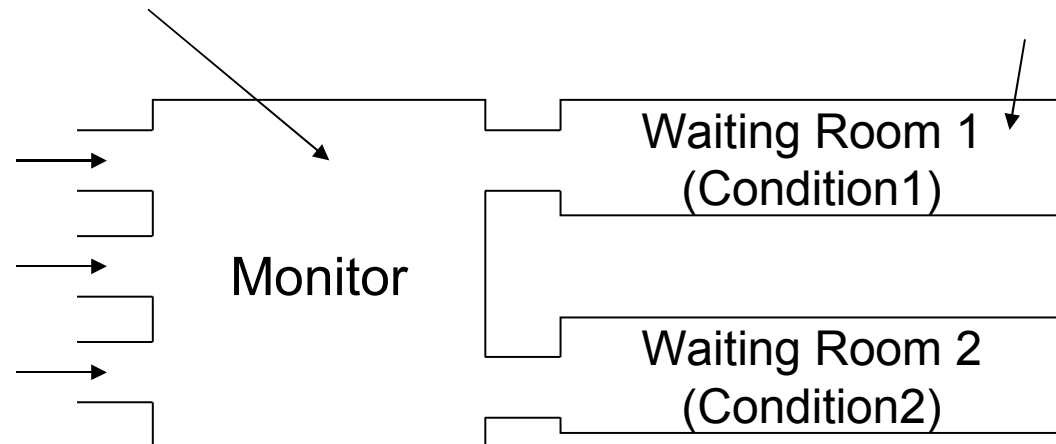
- In most programming languages (java, c++) using D.C.L optimization may result in a crash
- Correct code is the previous one:

// correct code

```
class Foo {  
    private Helper = null;  
    public synchronized Helper getHelper() {  
        if (helper == null)  
            helper = new Helper();  
        return helper;  
    }  
    // other functions and members...  
}
```


Monitor object

- **As presented in the previous lecture:**
 - a lock that synchronizes access to the object
 - Condition variables for waiting for specific events
 - Notify() wakes up waiting processes



Thread Pool Pattern

- **Steady stream of tasks to be performed, e.g.**

- A web server handling many clients
- A database server


- **Single thread:**

- Not optimal on multiCPU systems
- Asynchronous IO is difficult

- **Thread per request**

- Simplifies programming
- Creating a thread is an expensive operation

Avoid the expense of creating a new thread by reusing existing ones.



- A thread that finished a task takes another one from the queue.
- If there are no more tasks the thread may terminate or sleep.

Thread Pool in Java

- **ThreadPoolExecutor interface**
 - execute(Runnable command)
- **Decouples submitting a task and running it**
- **Several executors Possible:**
 - Direct – just starts the task (synchronously)
 - Thread per task – creates a new thread for each task
 - ThreadPoolExecutor
 - corePoolSize – minimum # of threads
 - maximumPoolSize – maximum # of threads
 - keepAliveTime – threads wait at most this amount of time


```
Executor executor = anExecutor;  
executor.execute(new RunnableTask1());  
executor.execute(new RunnableTask2());
```

```
class DirectExecutor  
    implements Executor {  
    public void execute(Runnable r) {  
        r.run();  
    }  
}
```

```
class ThreadPerTaskExecutor  
    implements Executor {  
    public void execute(Runnable r) {  
        new Thread(r).start();  
    }  
}
```

Multithreaded server


- **Pool of threads**
- **Each new request is taken by an available thread**
- **A possible solution:**
 - One thread reads request from the network (UNIX *select* call)
 - Creates a request and passes it to a working thread
- **Problems:**
 - Passing request between threads requires frequent context switching and synchronization
- **Can we do it better?**



Leader-Followers pattern

- Threads take turns in reading request
- After reading a request:
 - Wake up another thread to read the next request
 - Process the request
 - If there is no leader
 - Become the new leader
 - Else
 - wait in the pool

```
loop {  
    ReadARequest();  
    ThreadPool.Notify();  
    ProcessTheRequest();  
    if (there is another leader)  
        ThreadPool.Wait();  
}
```



Global values

- **The C errno value**
 - Global for the whole application
 - Is set after a system function's call has failed
- **When there are threads the value might get overwritten in another thread**
- **Possible solution: protect the value with a lock**
 - System function will acquire a lock when an error appears
 - User application will read the value and release the lock
- **Very bad a solution**
 - It is difficult to change system functions
 - A programmer may forget to release the lock – deadlock

Such a problem appears in other legacy systems as well.

Thread-specific storage

- Introduce a global access point for such a global variable
- But keep separate value that variable per thread
- Values can be associated with threads
 - Java: Thread class:
 - `public static Thread currentThread()`
 - When subclassed values can be kept in the Thread object
 - ThreadLocal object – a value, different for each thread
 - C (pthread):
 - `pthread_t pthread_self(void);`
 - Gives the thread's ID. ID's can be mapped to objects.
 - pthread also has functions for associating values with keys for a thread

```
#define errno (*(__errno()))
```


```
int *__errno() {  
    // find the pointer p to thread-specific errno  
    return p;  
}
```

Future Object

- **Some operations have to be processes asynchronously**
 - Using a Thread Pool
 - Via the operating system (async. IO)
- **Future Object encapsulates**
 - The asynchronous operation's code
 - Completion (called after the operation has been finished)
- **All code at the same place**
- **Example: web server serving a client**
 - Asynchronous operation: reading file from disk
 - Completion: Invoke another async. operation – send the file.
 - Another completion: e.g. close the connection

Future Object cont.

- At some point one has to use the result from the asynchronous IO.
- How to check if the task is finished?
 - call isDone() method of the Future
 - call get() method (waits for the result)
 - Use a queue of finished Futures
- Queue
 - The future at the end of its asynchronous operation puts its identifier into a queue
 - The main task waits on the queue for finished tasks

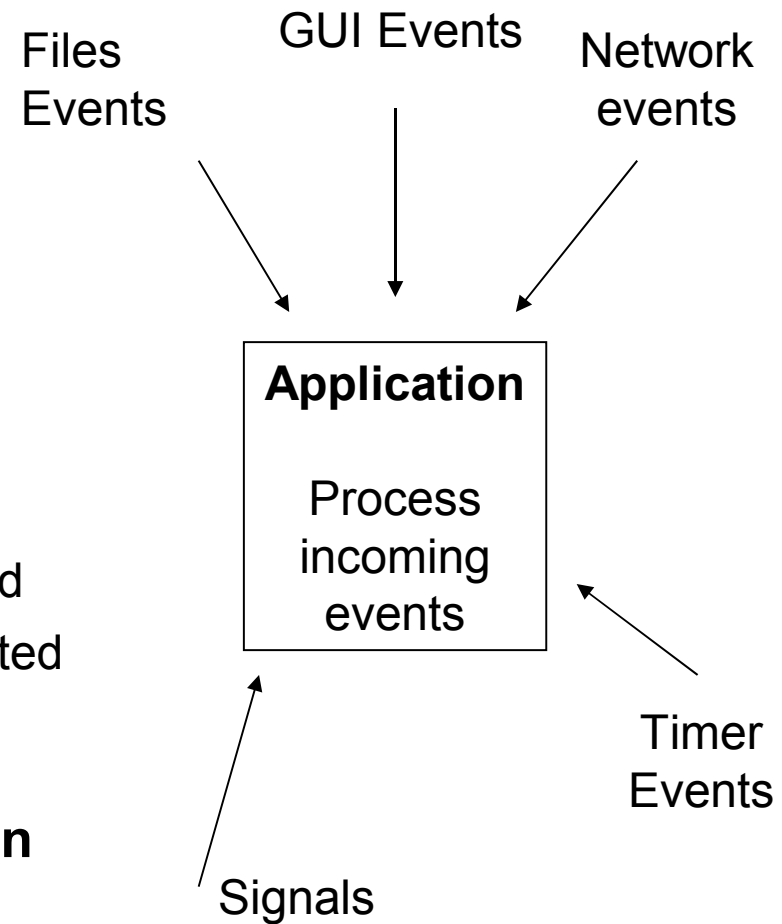


Futures in Java

- **Interface Future<V>:**
 - **boolean** cancel()
 - V get()
 - **boolean** isDone()
 - **boolean** isCancelled()
 - Additional methods may implement the completion
- **Class FutureTask<V>; additional methods:**
 - run() – the async. operation itself
 - done() – called when the state turns to *done*. May be overridden e.g. to notify the main thread
 - Additional methods for completion may be added when subclassed
- **FutureTask may be executed using Executors (e.g. ThreadPoolExecutor)**

Events

- **External „message”**
 - GUI event
 - The user clicked on a button
 - The user moved the mouse pointer
 - Network event
 - New network connection is coming
 - New data is arriving
 - More data can be sent right now
 - Operating system's events
 - Application is going to be terminated
 - Asynchronous IO has been completed
- **Or internal message, from another thread/process within the application**



Event Handles

- **Each event has an associated Event Handle**
 - An identifier of the event source (network connection, open file, GUI widget, etc.)
 - Usually provided by the O.S
- **Examples:**
 - For a network connection, files: descriptor number
 - For GUI: widgetID
 - Timers: timerID
- **There may be several kinds of events for one handle**
 - GUI widget:
 - MouseClicked
 - MouseMoved
 - For a file descriptor
 - Read event – more data to read
 - Write event – more data to write

Event-driven programming

- **Normal model of programming:**

```
Initialize();  
DoComputations();  
Terminate();
```

- **Event-driven programming**

- Events are processed in a loop (so called „Event Loop”)
- When no events no computations are performed

```
Initialize();  
while ((e = GetEvent()) != 0) {  
    HandleEvent(e);  
}  
Terminate();
```

Event Loop

- **The Event Loop is often provided by the O.S.**
- **Event handler – user defined function that handles an event**
 - Has to be registered for a specific event type
 - The O.S. will call back the event handler when a specific event appears
- **Inversed program's logic**
 - No main loop
 - Just event handlers should be written (+ the main initialization function)

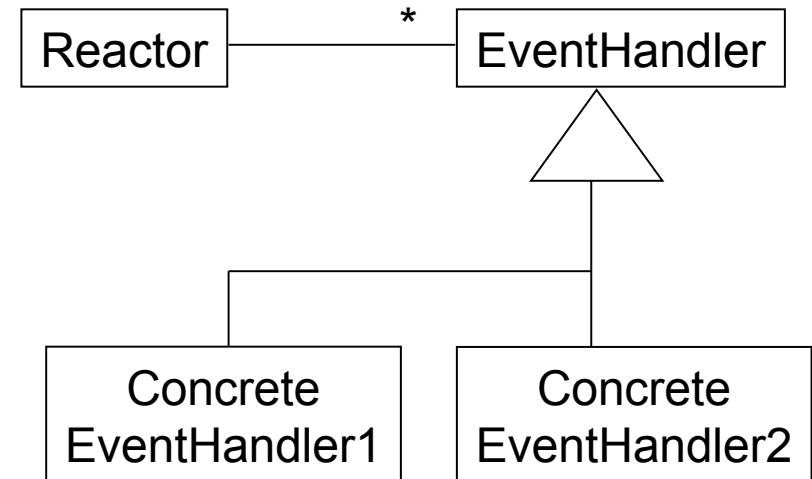
```
void SocketHandler() {  
    // read the incoming data  
}  
void ExitHandler() {  
    OS.StopEventLoop();  
}  
void main() {  
    socket s = OpenConnection();  
    OS.registerEventHandler(  
        EXIT_EVENT, ExitHandler);  
  
    OS.registerEventHandler(  
        SOCKET_EVENT, s,  
        SocketHandler);  
  
    OS.EventLoop();  
}
```

callback functions (event handlers)

Reactor pattern

Object-oriented way of implementing the event loop

- **EventHandler – abstract class**
 - Owns a handle
 - HandleEvent() method
- **Concrete Event Handlers inherit from EventHandler class**
- **Reactor**
 - EventLoop() ←
 - RegisterHandler(handler)
 - RemoveHandler(handler)



Loop:


- Get handles from all registered handlers
- Wait for an event from one of the handles (using the O.S)
- Call EventHandler->HandleEvent() for the handle returned by the demultiplexer

Java AWT – Event Handling

- **The Event Loop runs in a separate thread**
- **EventHandlers: `EventListener` interface**
 - Separate interfaces for different events
 - *MouseListener* for mouse events
 - *KeyListener* for keyboard events
- **Registering handlers to widgets:**
 - `aWidget.addMouseListener(aMouseListener);`
 - `aWidget.addKeyListener(aKeyListener);`
- **Handling events:**
 - Appropriate method of a listener is called:
 - e.g. `mousePressed()`, `mouseEntered()`
 - Each method takes a parameter: `EventObject`
 - `MouseEvent` for mouse events (mouse position, button state, etc.)

Handling asynchronous operations

- **Reactor Pattern can be used for handling asynchronous operations**
- **Remember the Future pattern?**



Summary

- **Design Patterns are successful solutions to common design problems**
- **There are useful patterns for concurrent programs**
- **Synchronization patterns: Scoped Locking, Thread-Safe Interface, Synchronizing Proxy, Read-Write Lock, Monitor**
- **Concurrency Patterns: Thread Pool, Leaders-Followers, Thread-Specific Storage**
- **Event-Handling Patterns: Future, Reactor, Proactor**

Further Reading and materials used

- Schmidt, Stal, Rohnert, Buchmann – *Pattern-Oriented Software Architecture, Patterns for Concurrent and Networked Objects, Volume 2*
- Gamma, Helm, Johnson, Vlissides – *Design Patterns: Elements of Reusable Object-Oriented Software*
- Java standard library
- Wikipedia
- Contact: mbiskup@mimuw.edu.pl